

IBM Cloud Containers Workshop

Part2 - Orchestration



Introduction to Orchestration

- Container Orchestration = Scheduling + Cluster management + Discovery

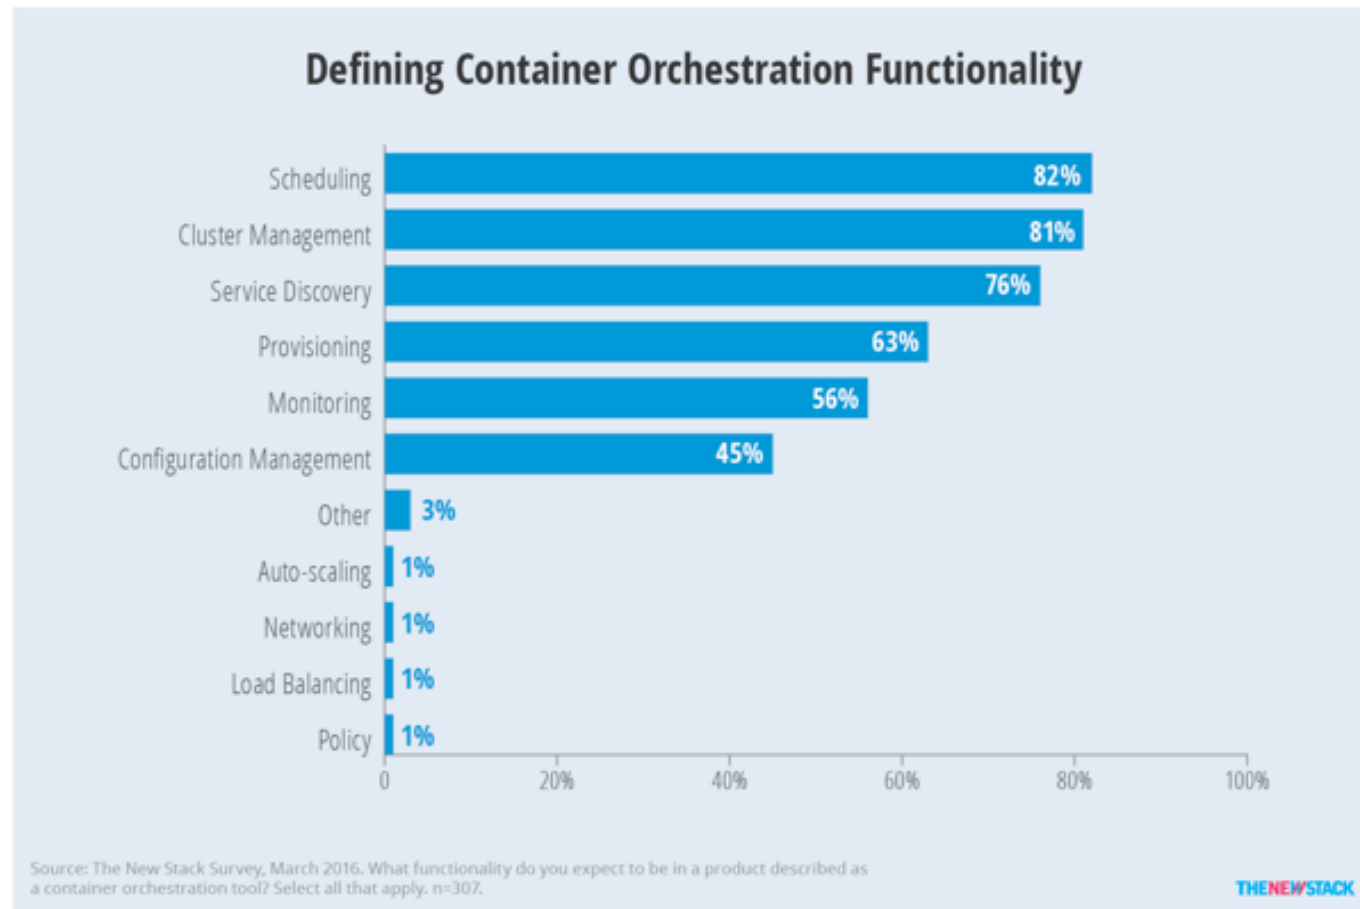
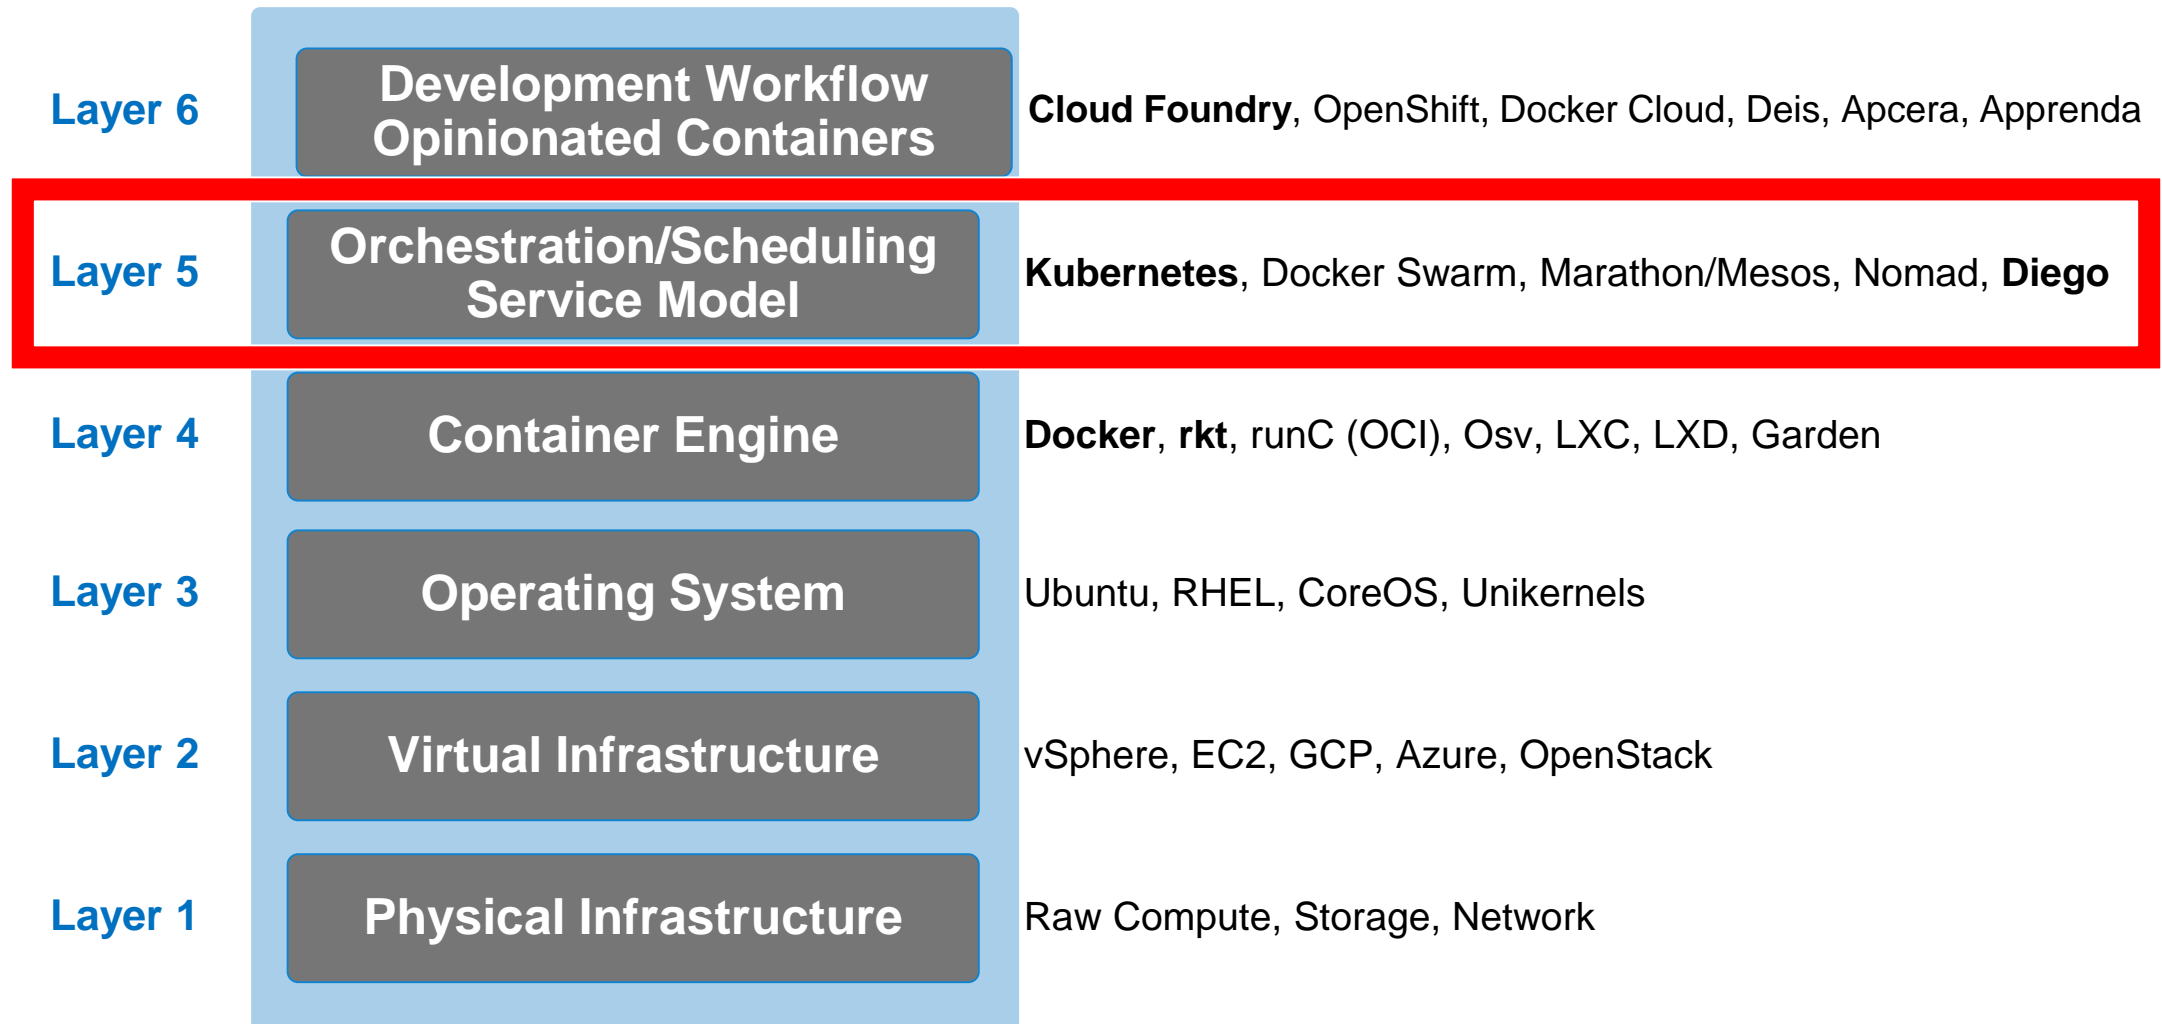
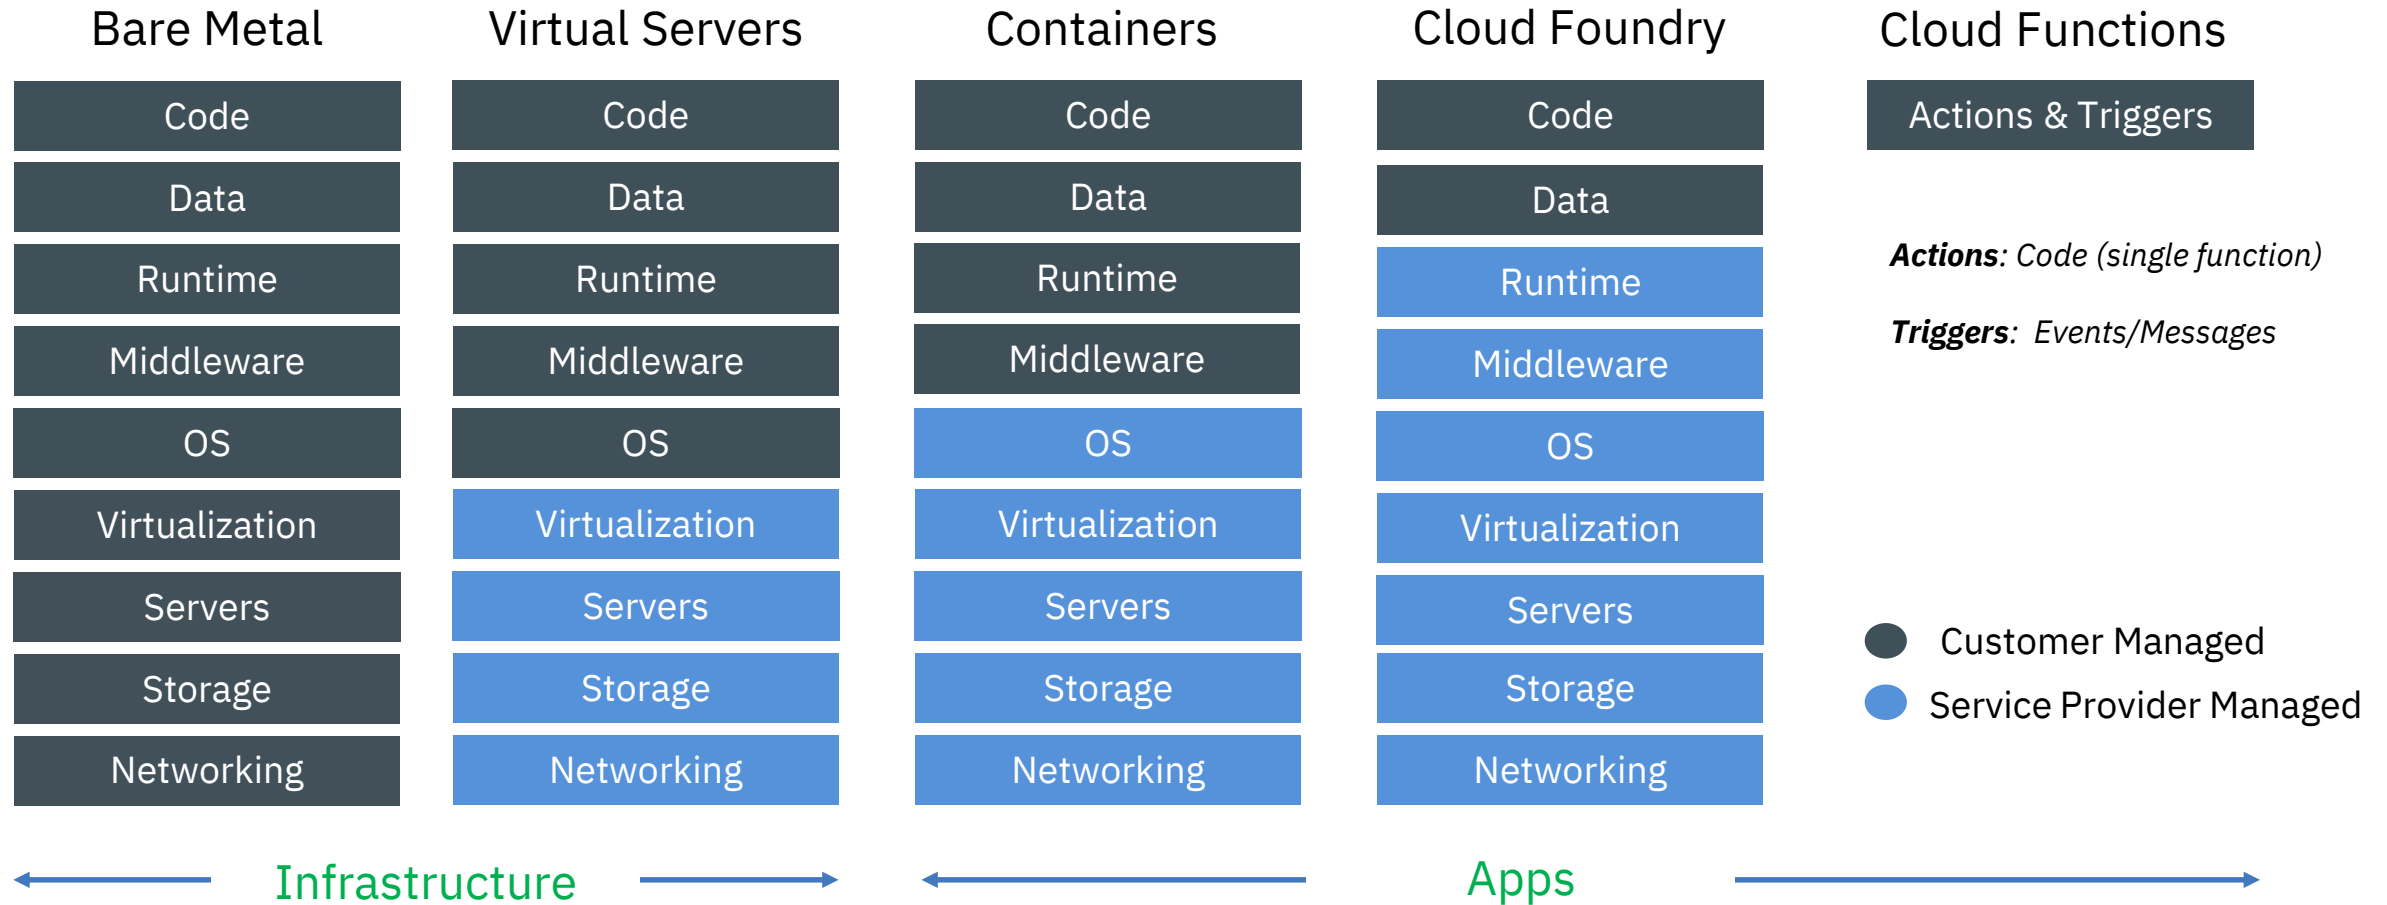


Figure 3: Only 45 percent of respondents consider configuration management to be part of a container orchestration product.

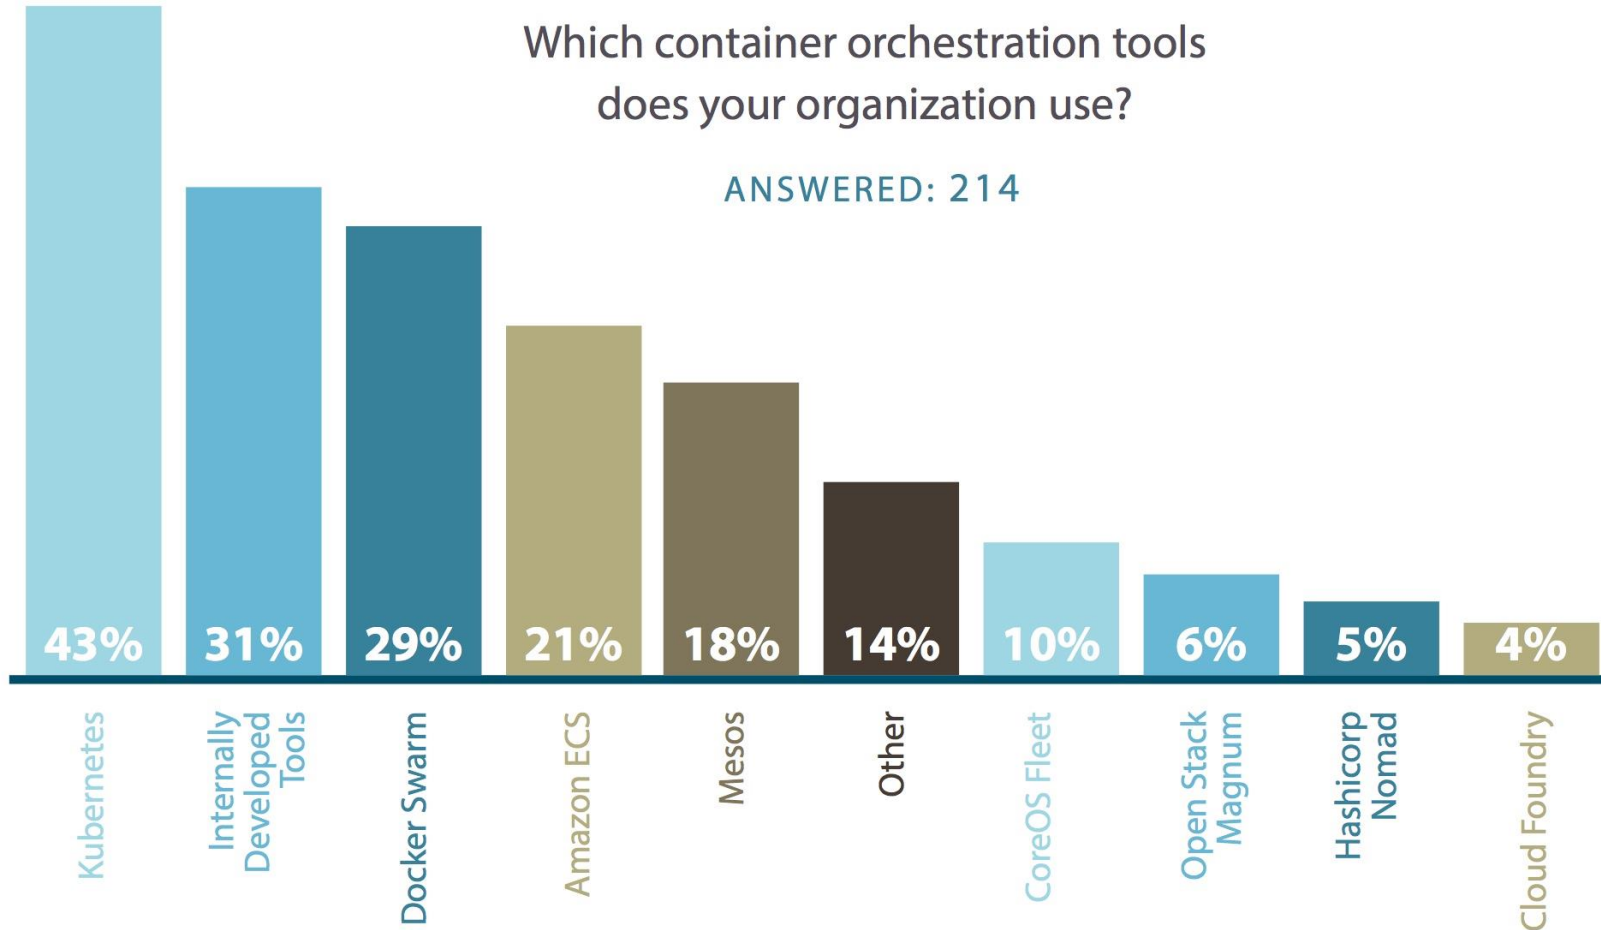
Container Orchestration with Kubernetes



• Cloud Computing – Levels of Responsibility



Different Orchestration Tools

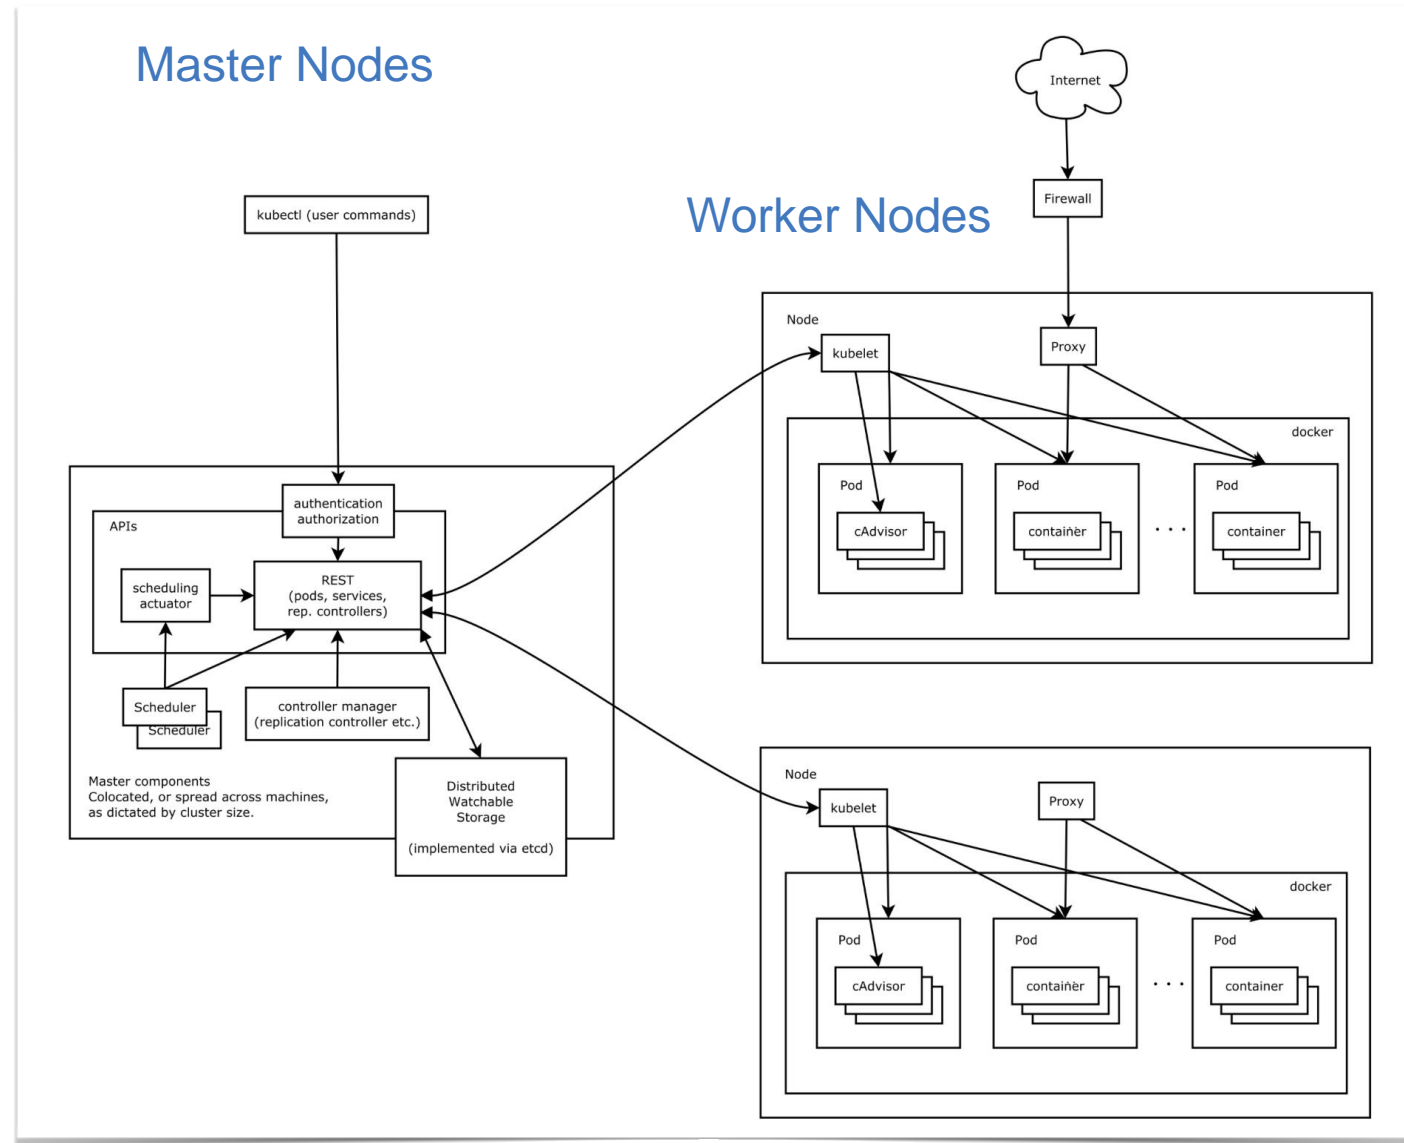


What is Kubernetes – K8S?

- Kubernetes is an [open-source platform for automating deployment, scaling, and operations of application containers](#) across clusters of hosts, providing container-centric infrastructure.
- Container orchestrator
- Runs and manages containers
- Supports **multiple cloud and bare-metal** environments
- Inspired and informed by Google's experiences and internal systems
- **100% Open source**, written in Go
- **Manage applications**, not machines
- Rich ecosystem of plug-ins for scheduling, storage, networking

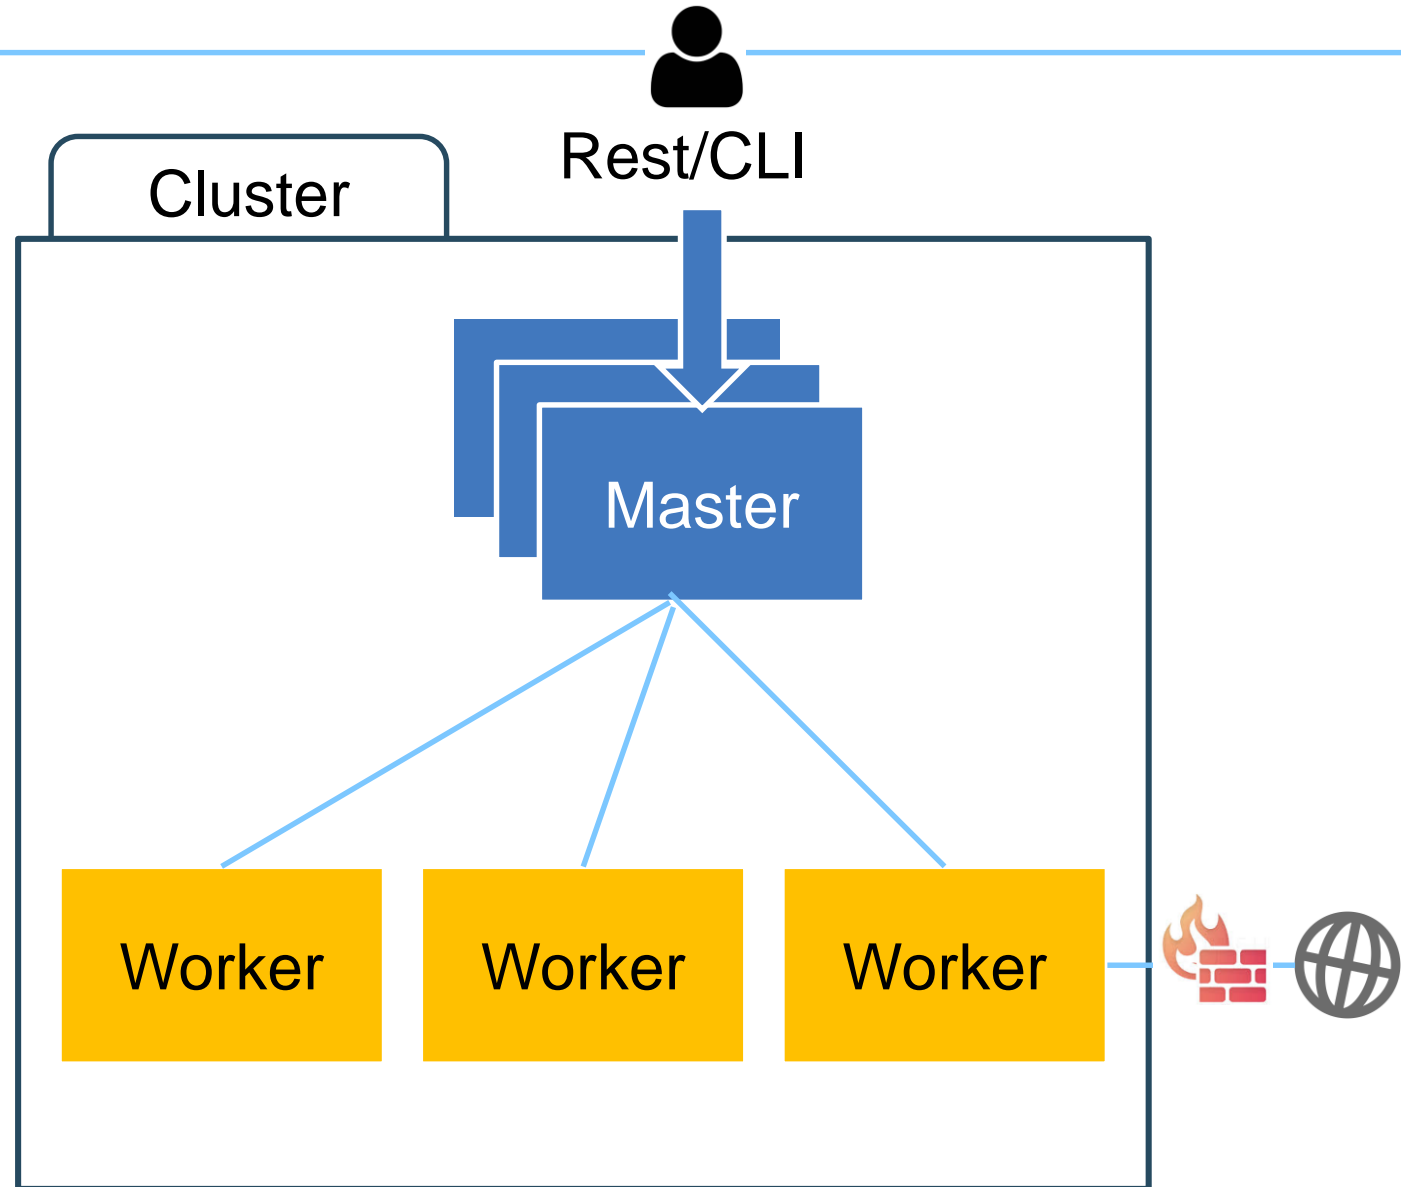


Kubernetes Overview



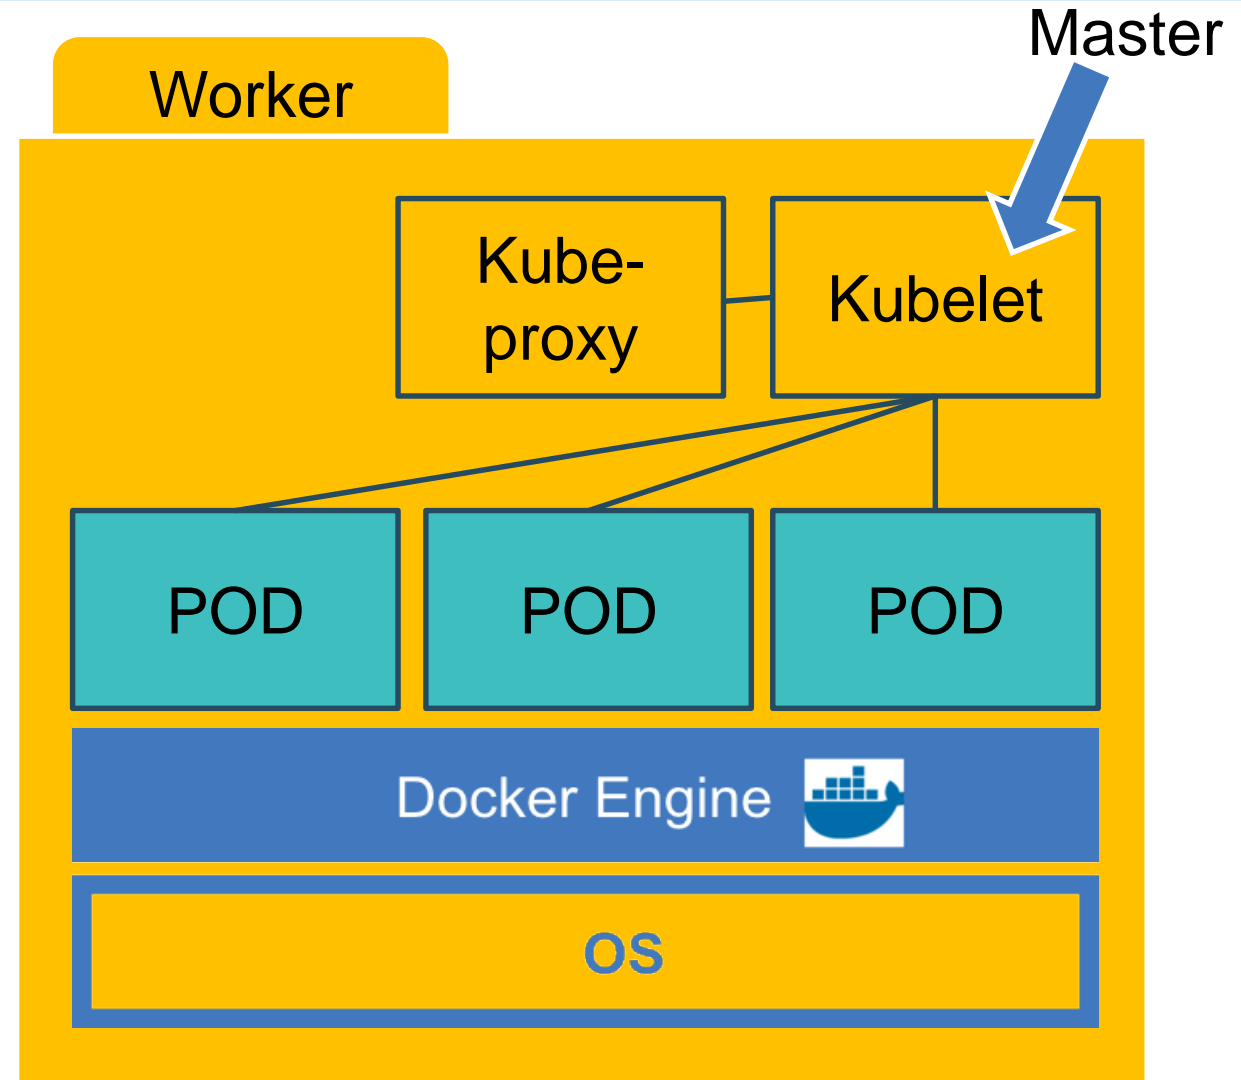
Cluster & Nodes

- Cluster is a set of nodes
- Nodes – hosts that run Kubernetes applications
- Master Node(s) – controls and manages the cluster :
 - Etcd
 - API Server
 - Controller Manager
 - Scheduler
- Worker Nodes – where the applications run



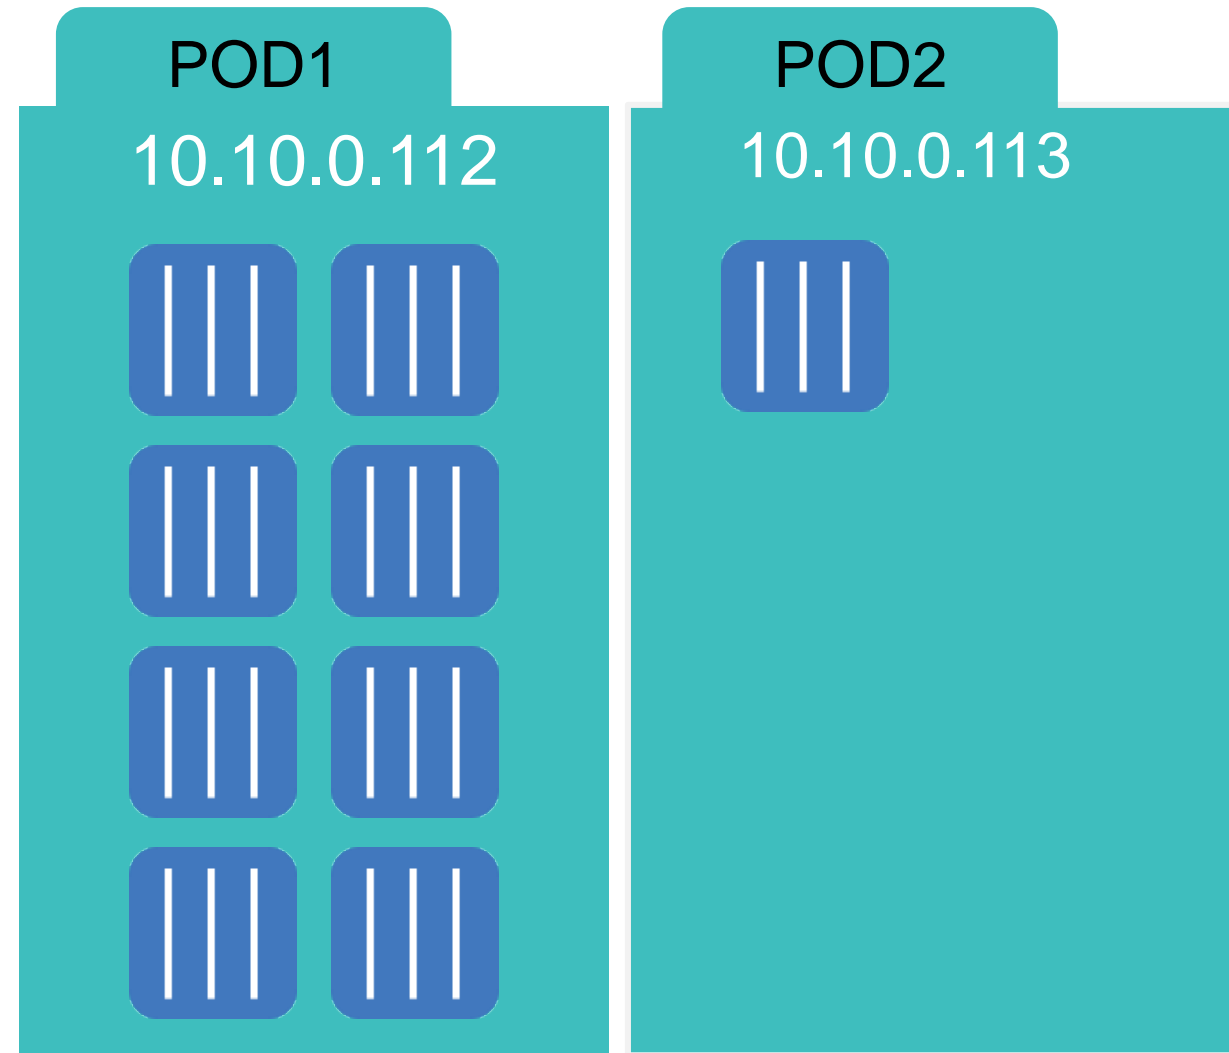
Worker Nodes

- This is where your applications are running in **PODs**
- Host Kubernetes services
- Runs the **kubelet** agent to control the node from the master
- **Kube-proxy** (network proxy service responsible for routing activities for inbound or ingress traffic to the PODs)
- Docker Engine Host



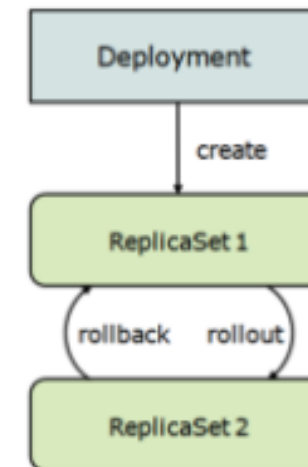
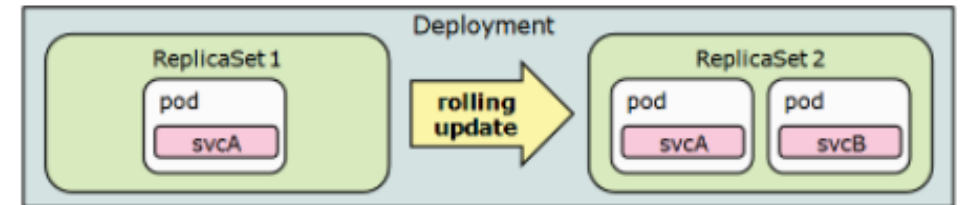
PODs

- Smallest **deployment unit** in K8s
- **Collection of containers** that run on a worker node
- Each has its **own IP**
- Pod shares a PID namespace, network, and hostname
- Inside a POD, from the network point of view, containers are in Network Container Mode (see Networking in Part #1)
- Important : check Network Ports



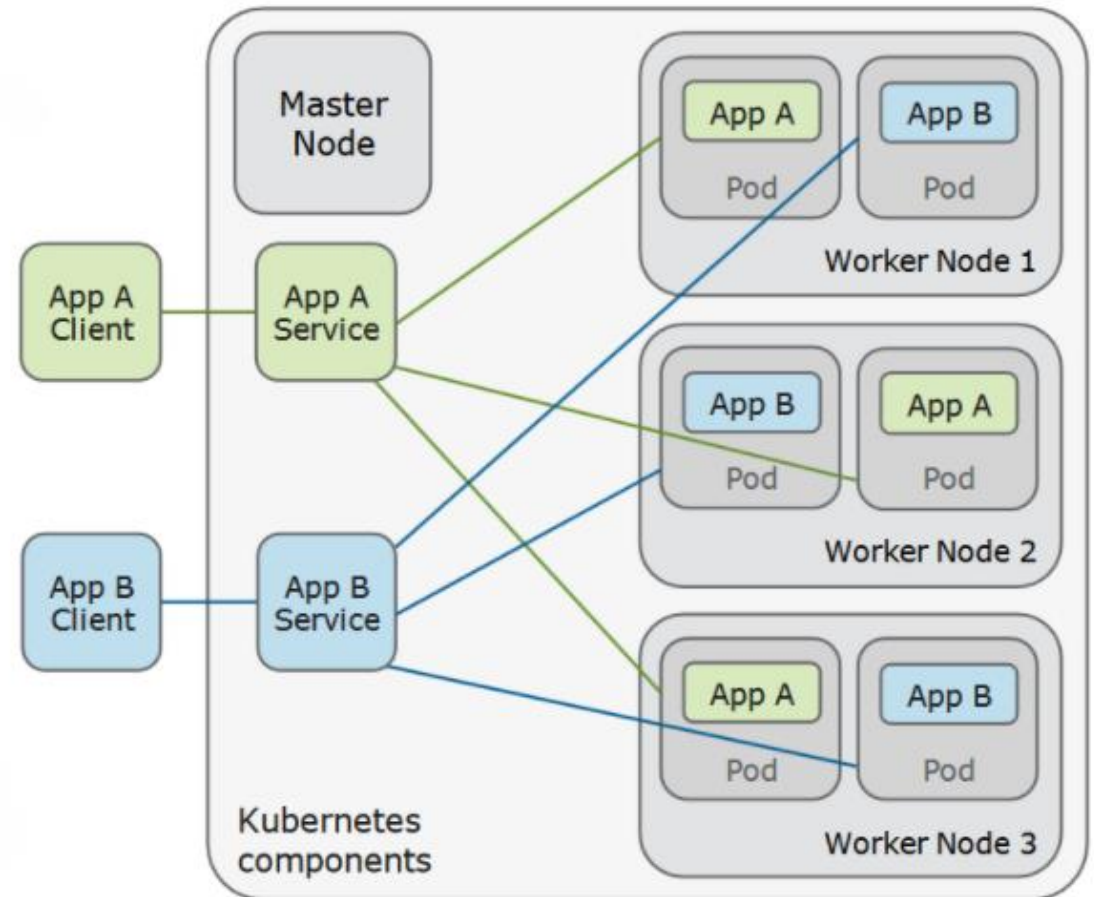
Controllers

- They control the **workloads**
- Different types : CronJob, DaemonSet, **Deployment**, Jobs, **Replicaset**, Replication-Controller
- A **Deployment** controller provides declarative updates for **Pods** and **ReplicaSets**.
- A ReplicaSet is a template that describes specifically what each pod should contain and it ensures that a specified number of pod replicas are running at any given time. It ensures **availability** and **scalability**



Kubernetes Services

- A service is a collection of pods exposed as an endpoint. The service propagates state and networking information to all worker nodes.
 - **ClusterIP**: exposes the cluster's internal IP (no communication with the outside)
 - **NodePort**: exposes the service on each node's IP at a static port
 - **LoadBalancer**: exposes the service externally by using a cloud provider's load balancer
 - **ExternalName**: maps the service to an external name, such as abc.xyz.example.com



Labels and Selectors

- Labels
 - **Metadata** assigned to K8s resources
 - Key-value pairs for identification
 - Critical to K8s as it relies on querying the cluster for resources that have certain labels
- **Selectors**
 - Equality (= or not =)
 - Set Based (in or not in)

```
"metadata": {  
  "labels": {  
    environmemnt: prod,  
    app: invoice  
  }  
}
```

```
$ kubectl get pods -l  
environment=prod,app=invoice
```

Names and Namespaces

- Each resource object by type has a unique **name**.
- To achieve resource isolation, each **namespace** is a virtual cluster within the physical cluster. Higher level resource objects are scoped within namespaces. Low level resources (nodes, persistent volumes and namespaces themselves) are not in namespaces. Within a namespace, names of resources must be unique, but not across namespaces. Namespaces can divide cluster resources.
- There are two initial namespaces:
- *default* - This is the namespace for objects with no other namespace.
- *kube-system* - The namespace for objects created by the Kubernetes system.

```
$ kubectl get namespaces
NAME                STATUS    AGE
default             Active    1d
kube-system         Active    1d
```

Configuration components

- **ConfigMaps**
 - ConfigMaps allow you to decouple configuration artifacts from image content to keep containerized applications portable.
- **Secrets:**
 - Sensitive info that containers need to consume
 - Encrypted in special volumes mounted automatically



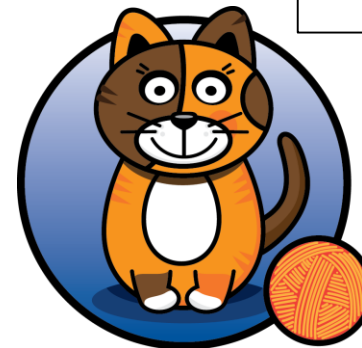
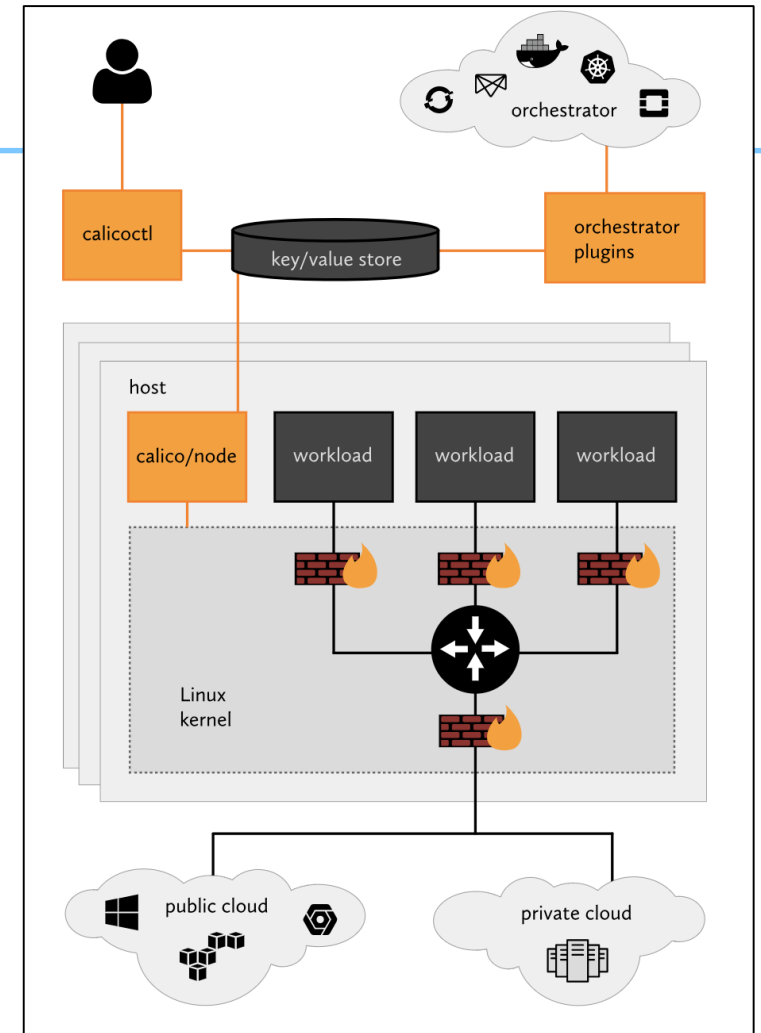
Configuration components

- **Network policy**
 - A Network policy is a specification of how groups of pods are allowed to communicate with each other and other network endpoints.
 - It uses labels to select pods and define rules which specify what traffic is allowed to the selected pods.



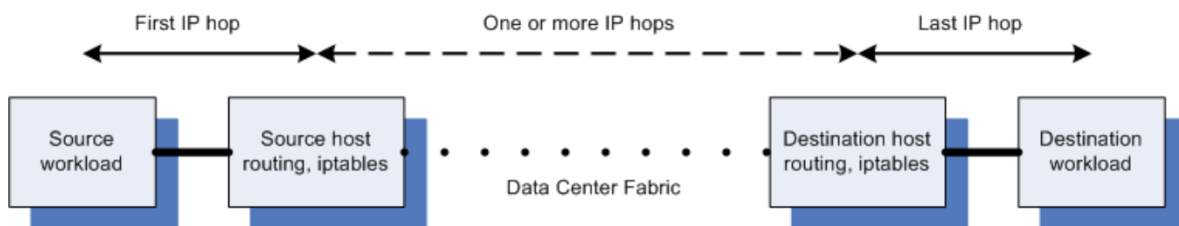
Calico

- A new approach to **virtual networking** and **network security** for **containers**, **VMs**, and **bare metal** services, that provides a rich set of security enforcement capabilities running on top of a highly scalable and efficient virtual network
- The **calico/node** Docker container runs on the Kubernetes master and each Kubernetes node in the cluster
- The **calico-cni plug-in** integrates directly with the Kubernetes kubelet process on each node to discover which pods have been created, and adds them to Calico networking
- The calico/kube-policy-controller container runs as a pod on top of Kubernetes and implements the NetworkPolicy API



PROJECT
CALICO

Calico Data Path: IP Routing and IPTABLES



Suppose that IPv4 addresses for the workloads are allocated from a datacenter-private subnet of 10.65/16, and that the hosts have IP addresses from 172.18.203/24. If you look at the routing table on a host you will see something like this:

```

ubuntu@calico-ci02:~$ route -n
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0          172.18.203.1   0.0.0.0         UG    0      0      0 eth0
10.65.0.0        0.0.0.0        255.255.0.0     U     0      0      0 ns-db03ab89-b4
10.65.0.21       172.18.203.126 255.255.255.255 UGH    0      0      0 eth0
10.65.0.22       172.18.203.129 255.255.255.255 UGH    0      0      0 eth0
10.65.0.23       172.18.203.129 255.255.255.255 UGH    0      0      0 eth0
10.65.0.24       0.0.0.0        255.255.255.255 UH     0      0      0 tapa429fb36-04
172.18.203.0     0.0.0.0        255.255.255.0   U     0      0      0 eth0
  
```

```

Chain cali-from-wl-dispatch (2 references)
target     prot opt source                destination
cali-from-wl-dispatch-0 all -- anywhere            anywhere    [goto] /* cali:eBnVcASLTvMFg9XV */
cali-from-wl-dispatch-1 all -- anywhere            anywhere    [goto] /* cali:MLdUGW-Du40orI9- */
cali-fw-cali2f9dbb62ca9 all -- anywhere            anywhere    [goto] /* cali:3L2Z8Gf72EsZIKAA */
cali-from-wl-dispatch-3 all -- anywhere            anywhere    [goto] /* cali:ZrLoQL8wJIHAgfVf */
cali-from-wl-dispatch-5 all -- anywhere            anywhere    [goto] /* cali:02H5UAoHug8r8x6m */
cali-from-wl-dispatch-6 all -- anywhere            anywhere    [goto] /* cali:s0d49VWBC4cE1fWn */
cali-from-wl-dispatch-9 all -- anywhere            anywhere    [goto] /* cali:X1yIW37zTvIP-0Y0 */
cali-fw-cali395ab80fdf all -- anywhere            anywhere    [goto] /* cali:ZeMp20hy0D-czLmE */
cali-fw-calic5f68922e7b all -- anywhere            anywhere    [goto] /* cali:w1Nq3uLVtq2qH3oD */
cali-from-wl-dispatch-d all -- anywhere            anywhere    [goto] /* cali:BhF3uys5r8zccmqw */
cali-from-wl-dispatch-e all -- anywhere            anywhere    [goto] /* cali:c0Xc1KJyb5GH35sQ */
cali-from-wl-dispatch-f all -- anywhere            anywhere    [goto] /* cali:c1Sy0ys7S0_GB2tR */
DROP       all -- anywhere            anywhere    /* cali:2P0cqaIbivIs3yNs */ /* Unknown interface */

Chain cali-from-wl-dispatch-0 (1 references)
target     prot opt source                destination
cali-fw-cali01b1faa5043 all -- anywhere            anywhere    [goto] /* cali:GGK4M4rUyo36iz8Yb */
cali-fw-cali092febb2d14 all -- anywhere            anywhere    [goto] /* cali:Pk1v50zerJBINX0F */
DROP       all -- anywhere            anywhere    /* cali:tRputS56cVeoNqYP */ /* Unknown interface */

Chain cali-from-wl-dispatch-1 (1 references)
target     prot opt source                destination
cali-fw-cali16dfc76a654 all -- anywhere            anywhere    [goto] /* cali:ymx3z1CMx7bhjz2E */
cali-fw-cali176404e5246 all -- anywhere            anywhere    [goto] /* cali:S1QW0TxEu1gsFIUc */
cali-fw-cali18ebc6b5707 all -- anywhere            anywhere    [goto] /* cali:tDxvYjyeEKdRtpvG */
cali-fw-cali1fd6184583a all -- anywhere            anywhere    [goto] /* cali:MgPmzB7fg0raq1Si */
DROP       all -- anywhere            anywhere    /* cali:cL5dCkd45Zcm9y7a */ /* Unknown interface */

Chain cali-from-wl-dispatch-3 (1 references)
target     prot opt source                destination
cali-fw-cali303beeb22d4 all -- anywhere            anywhere    [goto] /* cali:NC5e0XLCshu3R7VN */
cali-fw-cali30c3290b8cd all -- anywhere            anywhere    [goto] /* cali:UnaNmKqQW14RETt */
cali-fw-cali3d1c4438cf0 all -- anywhere            anywhere    [goto] /* cali:4g5gp_NGwTnnYHFV */
cali-fw-cali3e3e8a47393 all -- anywhere            anywhere    [goto] /* cali:CD1myuMSjL0P5ZG0 */
DROP       all -- anywhere            anywhere    /* cali:Inx8U2DMdYmYo0Y3 */ /* Unknown interface */
  
```

Service discovery

- Services need to discover each other dynamically, to get IP address and port detail to communicate with other services in the cluster
- Kubernetes provides two options for internal service discovery :
 - – **Environment variable**: When a new Pod is created, environment variables from older services can be imported. This allows services to talk to each other. This approach enforces ordering in service creation.
 - – **DNS**: Every service registers to the DNS service; using this, new services can find and talk to other services. Kubernetes provides the **kube-dns** service for this.

Autoscaling

- Kubernetes implements autoscaling through Horizontal Pod Autoscaling (HPA). HPA automatically scales the number of pods in a replication controller, deployment, or replica set by matching the observed average CPU utilization to a specified target. HPA can also autoscale based on application-provided metrics.
- Metrics to drive this are fetched in two ways: direct Heapster access and REST client access. Kubernetes Heapster enables container cluster monitoring and performance analysis.. The default configuration is to:
 - query every 30 seconds
 - maintain 10% tolerance
 - wait 3 minutes after scale-up
 - wait another 5 minutes after scale-down
- The following kubectl command creates an HPA instance that maintains between 1 and 10 replicas of the pod controlled by the deployment. Additionally, the command maintains an average CPU utilization across all pods of 50%.

```
$ kubectl autoscale deployment <deployment-name> --cpu-percent=50  
--min=1 --max=10 deployment "<hpa-name>" autoscaled
```

kubectl Commands

```
kubectl run hello-world-deployment --image=registry.ng.bluemix.net/production/hello-world:2
```

```
philmacbook:Stage1 phil$ kubectl
kubectl controls the Kubernetes cluster manager.

Find more information at https://github.com/kubernetes/kubernetes.

Basic Commands (Beginner):
  create      Create a resource by filename or stdin
  expose      Take a replication controller, service, deployment or pod and expose it as a new Kubernetes Service
  run         Run a particular image on the cluster
  set         Set specific features on objects

Basic Commands (Intermediate):
  get         Display one or many resources
  explain     Documentation of resources
  edit        Edit a resource on the server
  delete      Delete resources by filenames, stdin, resources and names, or by resources and label selector

Deploy Commands:
  rollout     Manage a deployment rollout
  rolling-update Perform a rolling update of the given ReplicationController
  scale       Set a new size for a Deployment, ReplicaSet, Replication Controller, or Job
  autoscale   Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:
  certificate  Modify certificate resources.
  cluster-info Display cluster info
  top          Display Resource (CPU/Memory/Storage) usage.
  cordon      Mark node as unschedulable
  uncordon    Mark node as schedulable
  drain       Drain node in preparation for maintenance
  taint        Update the taints on one or more nodes

Troubleshooting and Debugging Commands:
  describe    Show details of a specific resource or group of resources
  logs        Print the logs for a container in a pod
  attach      Attach to a running container
```


(a few) kubectl commands

- CLUSTER :
 - Get the state of your cluster
 - `$ kubectl cluster-info`
 - Get all the nodes of your cluster
 - `$ kubectl get nodes`
- DEPLOYMENTS :
 - Get deployments from namespace «kube-system »
 - `$ kubectl get deployments -n kube-system`
 - Get details from deployment
 - `$ kubectl describe deploy <NAME_OF_DEPLOYMENT>`
- SERVICES :
 - Get info about the services of your cluster
 - `$ kubectl get services`
 - Get full config info about a Service
 - `$ kubectl get service <NAME_OF_SERVICE> -o json`
 - Delete a Service
 - `$ kubectl delete service NAME_OF_THE_SERVICE`
- PODS :
 - Get info about the pods of your cluster
 - `$ kubectl get pods --all-namespaces`
 - Get the IP of a Pod
 - `$ kubectl get pod <NAME_OF_POD> -template={{.status.podIP}}`
 - Delete a Pod
 - `$ kubectl delete pod NAME`

kubectl apply -f filename.yml

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: hw-demo-deployment
spec:
  replicas: 3
  template:
    metadata:
      name: pod-liveness-http
    labels:
      run: hw-demo-health
      test: hello-world-demo
```

```
spec:
  containers:
    - name: hw-demo-container
      image: "registry.ng.bluemix.net/pod1/hello-world:2"
      imagePullPolicy: Always
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
        initialDelaySeconds: 5
        periodSeconds: 5
```

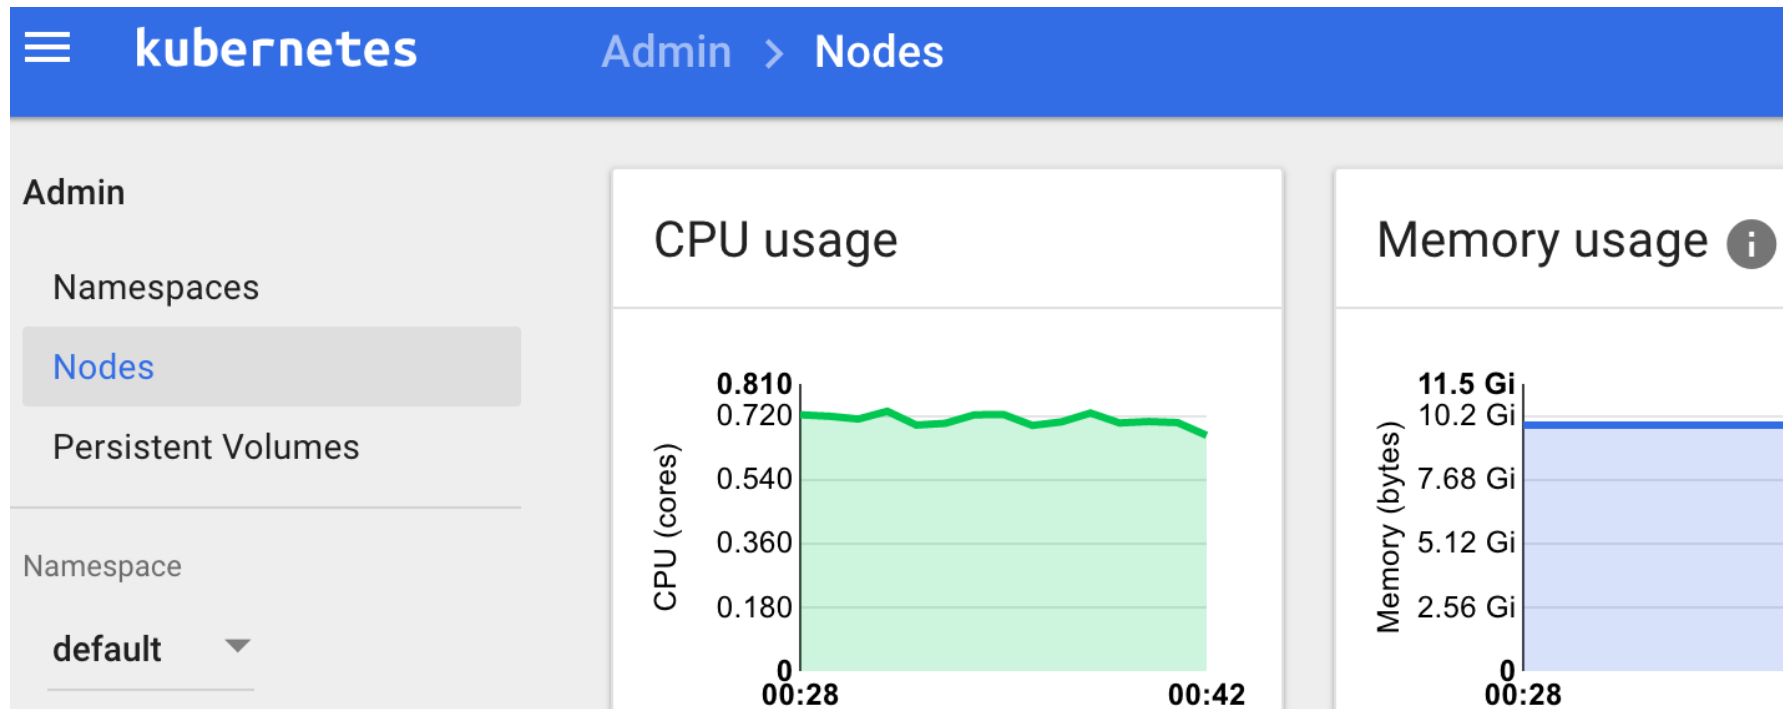
kubectl delete -f filename.yml

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: hw-demo-deployment
spec:
  replicas: 3
  template:
    metadata:
      name: pod-liveness-http
    labels:
      run: hw-demo-health
      test: hello-world-demo
```

```
spec:
  containers:
    - name: hw-demo-container
      image: "registry.ng.bluemix.net/pod1/hello-world:2"
      imagePullPolicy: Always
      livenessProbe:
        httpGet:
          path: /healthz
          port: 8080
        initialDelaySeconds: 5
        periodSeconds: 5
```

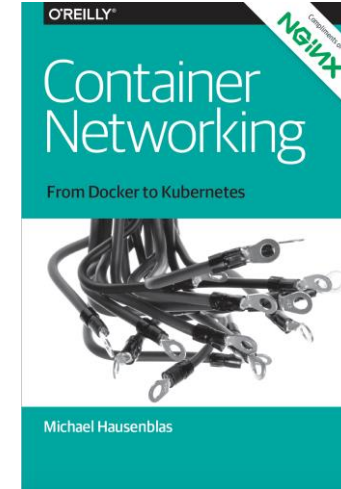
Monitoring Containers

- **Integrated** logging and monitoring on IBM Cloud based on ELK stack
- Native Kubernetes **dashboard** or API



Books, eBooks and links

- Getting Started with Kubernetes - Second Edition
- Mastering Kubernetes
- Container Networking



- <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
- <https://kubernetes.io/docs/concepts/>





Intelligent Scheduling



Automated rollouts and rollbacks



Design Your Own Cluster



Container Security & Privacy



Service discovery & load balancing



Secret & configuration management



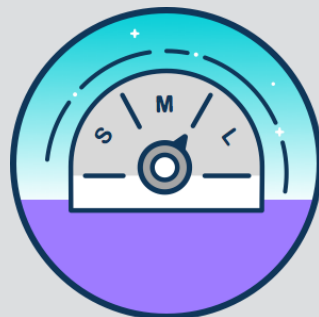
Simplified Cluster Management



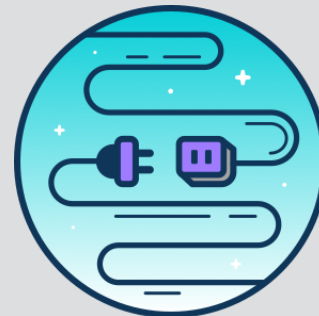
Native Kubernetes Experience



Self-healing



Horizontal scaling



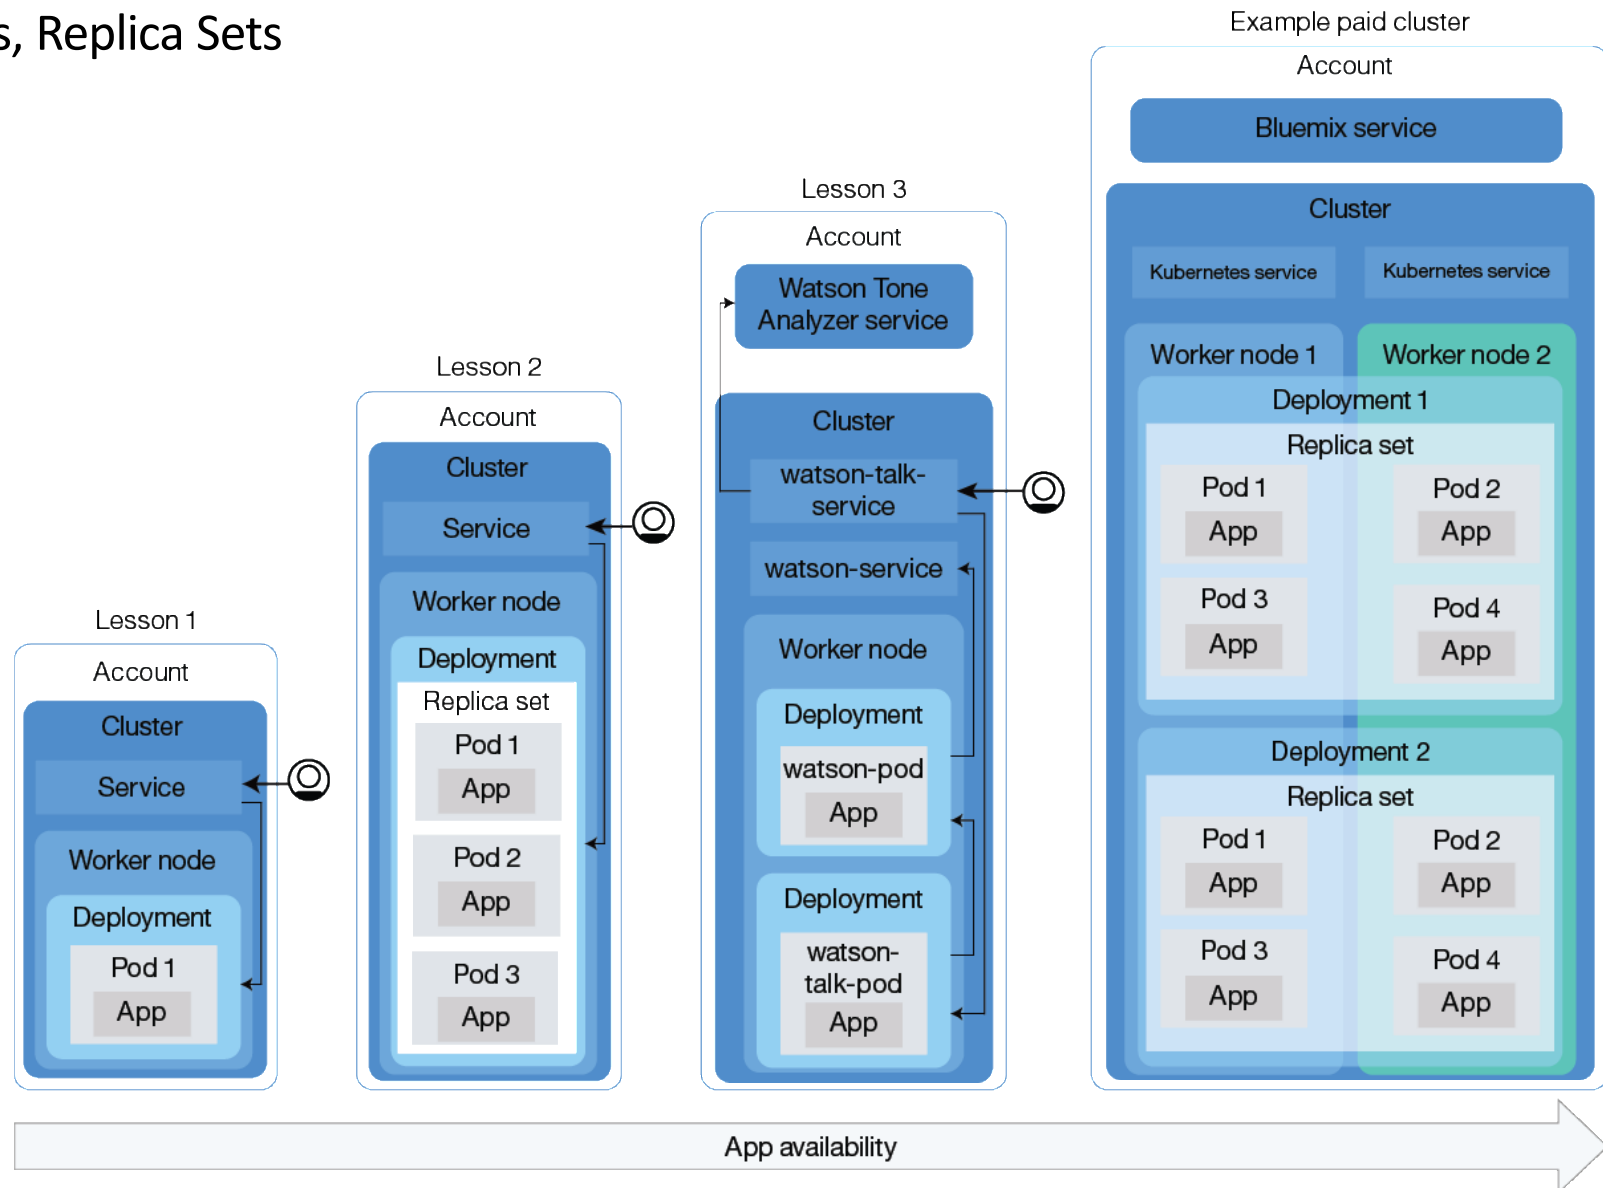
Leverages IBM Cloud & Watson



Integrated Operational Tools

Container Application Resiliency – High Availability Patterns

Leveraging container Pods, Replica Sets and Worker Nodes



Example : Docker Build

```
root:[Lab 1]: docker build -t mycluster.icp:8500/default/hello-world .
Sending build context to Docker daemon 15.36kB
Step 1/6 : FROM node:9.4.0-alpine
9.4.0-alpine: Pulling from library/node
605ce1bd3f31: Pull complete
fe58b30348fe: Pull complete
46ef8987ccbd: Pull complete
Digest: sha256:9cd67a00ed111285460a83847720132204185e9321ec35dacec0d8b9bf674adf
Status: Downloaded newer image for node:9.4.0-alpine
----> b5f94997f35f
Step 2/6 : COPY app.js .
----> bbfe2d4ee8d0
Step 3/6 : COPY package.json .
----> bf0e9fcc6637
Step 4/6 : RUN npm install --silent --omit=dev
----> 46e300000000
Step 5/6 : EXPOSE 8080
----> Running in 535c253950bc
----> 7f2e3656c237
Removing intermediate container 535c253950bc
Step 6/6 : CMD node app.js
----> Running in d8aea9eaaf3b
----> b5107d9859f5
Removing intermediate container d8aea9eaaf3b
Successfully built b5107d9859f5
Successfully tagged mycluster.icp:8500/default/hello-world:latest
```


Example : Docker Push

```
[root:[Lab 1]: docker login mycluster.icp:8500
Username: admin
Password:
Login Succeeded
[root:[Lab 1]:
[root:[Lab 1]:
[root:[Lab 1]: docker push mycluster:8500/default/hello-world
The push refers to a repository [mycluster:8500/default/hello-world]
An image does not exist locally with the tag: mycluster:8500/default/hello-world
[root:[Lab 1]: docker push mycluster:8500/default/hello-world:latest
The push refers to a repository [mycluster:8500/default/hello-world]
An image does not exist locally with the tag: mycluster:8500/default/hello-world
[root:[Lab 1]: docker push mycluster.icp:8500/default/hello-world:latest
The push refers to a repository [mycluster.icp:8500/default/hello-world]
3df45be69528: Pushed
11c7c90fdcae: Pushed
9dcd61b5afdc: Pushed
0804854a4553: Pushed
6bd4a62f5178: Pushed
9dfa40a0da3b: Pushed
latest: digest: sha256:f8216ed41187ef0a827de0caadbc0819e77e813220109ec78ae99049759eaca2 size: 1576
[root:[Lab 1]:
```

Example : Kubernetes Deployment, Expose, Describe

```
[philmacbook:Stage1 phil$ kubectl run hello-world-deployment --image=registry.ng.bluemix.net/prod1/hello-world:1
deployment "hello-world-deployment" created
```

```
[philmacbook:Stage1 phil$
```

```
[philmacbook:Stage1 phil$ kubectl expose deployment/hello-world-deployment --type=NodePort --port=8080 --name=hello-world-service
service "hello-world-service" exposed
philmacbook:Stage1 phil$
```

```
[root:[Lab 1]: kubectl describe service hello-world-service
Name:                hello-world-service
Namespace:           default
Labels:              run=hello-world-deployment
Annotations:         <none>
Selector:            run=hello-world-deployment
Type:               NodePort
IP:                 10.0.0.137
Port:               <unset> 8080/TCP
TargetPort:         8080/TCP
NodePort:           <unset> 30742/TCP
Endpoints:          10.1.210.151:8080
Session Affinity:    None
External Traffic Policy: Cluster
Events:             <none>
[root:[Lab 1]:
```

Example : Access to the app



Hello world from hello-world-deployment-69d8fcb8db-mbbw4! Your app is up and running in a cluster!